

How to Reliably Measure Software Performance

Augusto de Oliveira, Kemal Akkoyun

FOSDEM 2026

nature

News | Published: 22 September 2011


Particles break light-speed limit

[Geoff Brumfiel](#)

[Nature](#) (2011) | [Cite this article](#)


2879 Accesses | **13** Citations | **854** Altmetric | [Metrics](#)

CERN to Gran Sasso Neutrino Beam



732 km neutrino beam path from CERN in Geneva to Gran Sasso [1]

100 M€_[2]



Loose fiber optic cable that caused the measurement error [3]

How to:

1. Control your benchmarking environment
2. Design your benchmarks
3. Interpret benchmark results
4. Integrate benchmarks into your workflows

How to control your benchmarking environment

Layer	Noise Sources	Mitigations
External	Network Temperature Vibration	Use dedicated hardware
Application	Memory layout Compilation/linking	Set up fixed builds (e.g., disable ASLR)
Kernel	Filesystem cache Scheduling	Set CPU affinity Set process priority Warm up or drop caches
CPU	Simultaneous multithreading (SMT) contention Dynamic frequency scaling (DFS)	Disable SMT Disable DFS

Layer	Noise Sources	Mitigations
External	Network Temperature Vibration	Use dedicated hardware
Application	Memory layout Compilation/linking	Set up fixed builds (e.g., disable ASLR)
Kernel	Filesystem cache Scheduling	Set CPU affinity Set process priority Warm up or drop caches
CPU	Simultaneous multithreading (SMT) contention Dynamic frequency scaling (DFS)	Disable SMT Disable DFS

Bakhvalov, "Performance Analysis and Tuning on Modern CPUs", Appendix A [5]

```
# Set CPU affinity
taskset -c 0 ./benchmark

# Increase process priority
nice -n -5 ./benchmark

# Drop filesystem cache
echo 3 > /proc/sys/vm/drop_caches && sync


# Disable Simultaneous Multithreading (SMT)
echo off > /sys/devices/system/cpu/smt/control

# Disable Dynamic Frequency Scaling (DFS)
echo 1 > /sys/devices/system/cpu/intel_pstate/no_turbo


echo 2500000 > /sys/devices/system/cpu/cpu*/cpufreq/scaling_min_freq
echo 2500000 > /sys/devices/system/cpu/cpu*/cpufreq/scaling_max_freq

echo performance > /sys/devices/system/cpu/cpu*/cpufreq/scaling_governor
```

What's the impact of disabling SMT?




SMT enabled: hardware threads compete for resources



SMT disabled: hardware thread has exclusive access to resources

What's the impact of disabling SMT?


m5.metal, clock rate pinned, scaling governor set to "performance"
2 CPU-intensive tasks on same core (smt) vs. separate cores (no-smt)



Thread	mean ± stddev	coeff. of variation
smt-1	1537.64 ± 367.29 ms	23.887 %
smt-2	1536.88 ± 366.84 ms	23.869 %
no-smt-1	737.37 ± 0.32 ms	0.044 %
no-smt-2	737.93 ± 1.74 ms	0.235 %

What's the impact of disabling SMT?


m5.metal, clock rate pinned, scaling governor set to "performance"
2 CPU-intensive tasks on same core (smt) vs. separate cores (no-smt)



Thread	mean ± stddev	coeff. of variation
smt-1	1537.64 ± 367.29 ms	23.887 %
smt-2	1536.88 ± 366.84 ms	23.869 %
no-smt-1	737.37 ± 0.32 ms	0.044 %
no-smt-2	737.93 ± 1.74 ms	0.235 %

100x less variation

What's the impact of disabling DFS?




DFS enabled: CPU frequency is automatically set by the Scaling Governor and the Scaling Driver

What's the impact of disabling DFS?

m5.metal, SMT disabled

Varying number of CPU-intensive tasks on the same core with DFS on vs. off




Thread	mean ± stddev	coeff. of variation
dfs-1	533.97 ± 2.046 ms	0.383 %
dfs-8	578.67 ± 0.287 ms	0.050 %
no-dfs-1	738.18 ± 0.306 ms	0.041 %
no-dfs-8	739.18 ± 0.351 ms	0.047 %

What's the impact of disabling DFS?

m5.metal, SMT disabled

Varying number of CPU-intensive tasks on the same core with DFS on vs. off



Thread	mean ± stddev	coeff. of variation
dfs-1	533.97 ± 2.046 ms	0.383 %
dfs-8	578.67 ± 0.287 ms	0.050 %
no-dfs-1	738.18 ± 0.306 ms	0.041 %
no-dfs-8	739.18 ± 0.351 ms	0.047 %

10x less variation

Removes dynamic frequency scaling as a source of noise

How to Design Benchmarks

"All happy families are alike; each unhappy family is unhappy in its own way."

— Leo Tolstoy, *Anna Karenina*

"All happy ~~families~~ are alike; each unhappy ~~family~~ is unhappy in its own way."
benchmarks **benchmark**

**representative and
repeatable**

An unhappy benchmark

- Goal: Measuring dd-trace-java instrumentation overhead on a Spring app.

An unhappy benchmark


- Goal: Measuring dd-trace-java instrumentation overhead on a Spring app.
- System under test: Spring Petclinic instrumented (or not) with dd-trace-java.

An unhappy benchmark


- Goal: Measuring dd-trace-java instrumentation overhead on a Spring app.
- System under test: Spring Petclinic instrumented (or not) with dd-trace-java.
- Workload: As many requests as possible by 5 concurrent users.

An unhappy benchmark

- Goal: Measuring dd-trace-java instrumentation overhead on a Spring app.
- System under test: Spring Petclinic instrumented (or not) with dd-trace-java.
- Workload: As many requests as possible by 5 concurrent users.
- 20 second warmup, 15 seconds of actual measurements.





Many **false positives**, unnacceptably **high CoV** (= standard deviation / mean) of 11.80%.




Many **false positives**, unnacceptably **high CoV** (= standard deviation / mean) of 11.80%.

Are we running the benchmark long enough?






Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).




Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

For how long should we run the benchmark?




# measurements	CoV
30	6.95%
60	5.23%
90	4.59%



# measurements	CoV
30	6.95%
60	5.23%
90	4.59%


Tip #2: Collect enough samples to reduce intra-run variation.




# measurements	CoV
30	6.95%
60	5.23%
90	4.59%


Tip #2: Collect enough samples to reduce intra-run variation.


But what about inter-run variation?




Impact of initial state on FFT benchmark results [4]







Run #	mean ± stddev	CoV
1	20.08 ± 0.63 ms	3.16%
2	20.63 ± 0.56 ms	2.72%
3	20.31 ± 0.45 ms	2.23%
4	20.19 ± 0.54 ms	2.66%
5	20.26 ± 0.63 ms	3.11%
all	20.29 ± 0.60 ms	2.94%



Run #	mean ± stddev	CoV
1	20.08 ± 0.63 ms	3.16%
2	20.63 ± 0.56 ms	2.72%
3	20.31 ± 0.45 ms	2.23%
4	20.19 ± 0.54 ms	2.66%
5	20.26 ± 0.63 ms	3.11%
all	20.29 ± 0.60 ms	2.94%

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation.

Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

Tip #2: Collect enough samples to reduce intra-run variation ($N \geq 30$).

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation ($M \geq 5$).

CoV: 11.80% → **2.94%**

Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

Tip #2: Collect enough samples to reduce intra-run variation ($N \geq 30$).

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation ($M \geq 5$).

CoV: 11.80% → **2.94%**

Tip #4: Use deterministic inputs.

Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

Tip #2: Collect enough samples to reduce intra-run variation ($N \geq 30$).

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation ($M \geq 5$).

CoV: 11.80% → **2.94%**

Tip #4: Use deterministic inputs.

Tip #5: Use load generators that avoid the coordinated omission problem (e.g., k6).

Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

Tip #2: Collect enough samples to reduce intra-run variation ($N \geq 30$).

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation ($M \geq 5$).

CoV: 11.80% → **2.94%**

Tip #4: Use deterministic inputs.

Tip #5: Use load generators that avoid the coordinated omission problem (e.g., k6).

But what about inter-experiment variation?

Tip #1: Run benchmarks for longer to uncover perturbations (e.g., warmup effects).

Tip #2: Collect enough samples to reduce intra-run variation ($N \geq 30$).

Tip #3: Rerun benchmarks multiple times to reduce inter-run variation ($M \geq 5$).

CoV: 11.80% → **2.94%**

Tip #4: Use deterministic inputs.

Tip #5: Use load generators that avoid the coordinated omission problem (e.g., k6).

But what about inter-experiment variation?

Tip #0: Control your benchmarking environment.

Interpreting Benchmark Results

Emphasize that simply comparing means or percentiles is not enough. Statistical methods such as t-tests are absolutely necessary to draw meaningful conclusions.

Shout out to Henrik Ingo's talk on changepoint detection.

Integrating Benchmarks Into Your Workflows

A series of screenshots showing the different ways in which we integrate benchmarks into our workflows at Datadog, including: a basic architecture slide, reporting capabilities, PR comments, performance quality gates, operational excellence reviews, etc.

Concluding slides

Summarize the takeaways.

References

- [1] Universität Münster. "Neutrino oscillations in the neutrino beam from CERN to Gran Sasso." <https://www.uni-muenster.de/Physik.KP/en/AGFrekers/forschung/opera.html>. Accessed Jan 2026.
- [2] CERN. (1999). "From Geneva to Gran Sasso in 2.5 milliseconds!". <https://home.cern/news/press-release/cern/geneva-gran-sasso-25-milliseconds>. Accessed Jan 2026.
- [3] Strassler, M. (2012). "OPERA: What Went Wrong." <https://profmattstrassler.com/articles-and-posts/particle-physics-basics/neutrinos/neutrinos-faster-than-light/opera-what-went-wrong/>. Accessed Jan 2026.
- [4] Kalibera, T., Bulej, L., and Tuma, P. (2005). "Benchmark Precision and Random Initial State." In *Proceedings of the International Symposium on Performance Evaluation of Computer and Telecommunication Systems (SPECTS)*, pages 182-196. SCS.
- [5] Bakhvalov, D. (2020). *Performance Analysis and Tuning on Modern CPUs*.
- [6] Valles, A. (2009). "Performance Insights to Intel Hyper-Threading Technology." <https://web.archive.org/web/20150217050949/https://software.intel.com/en-us/articles/performance-insights-to-intel-hyper-threading-technology/>. Accessed Jan 2026.
- [7] Gregg, B. (2014). "Frequency Trails: Outliers." <https://www.brendangregg.com/FrequencyTrails/outliers.html#Causes>. Accessed Jan 2026.
- [8] Gregg, B. (2020). "Systems Performance: Enterprise and the Cloud.", p. 233, "P-states and C-states."
- [9] Humenay, E., Tarjan, D., and Skadron, K. (2007). "Impact of Process Variations on Multicore Performance Symmetry."
- [10] Linux Kernel Documentation. "CPUFreq Governors." <https://www.kernel.org/doc/Documentation/cpu-freq/governors.txt>. Accessed Jan 2026.
- [11] ArchWiki. "CPU frequency scaling." https://wiki.archlinux.org/title/CPU_frequency_scaling. Accessed Jan 2026.
- [12] Intel. "Intel Server Board and System Products Update on Intel Turbo Boost Technology Support with Low Power Intel Xeon Processor 3400/5500/5600 Series." https://cdrdv2-public.intel.com/840590/white_paper_turbo_boost_on_low_power_processor.pdf. Accessed Jan 2026.